

Technical Report: Final Project DS 5110: Spotify Analytics Pipeline with AWS, Snowflake, and Power BI

Team Members: Ankur Haritosh, Pavan Kumar
Khoury College of Computer Sciences
Data Science Program

December 11, 2024

Contents

1	Introduction	2
2	Literature Review	2
3	Methodology	2
3.1	Data Collection	2
3.2	Data Preprocessing	2
3.3	Analysis Techniques	2
4	Results	2
5	Discussion	3
6	Conclusion	3
7	References	3
A	Appendix A: Code	3
B	Appendix B: Additional Figures	4
B.1	Pipeline Architecture	4
B.2	Data Model	5
B.3	Power BI Dashboards	6
B.4	Machine Learning Model	7
B.5	Activity Diagram	7
B.6	Class Diagram	8

1 Introduction

This project focuses on building a fully automated serverless ETL pipeline to extract, transform, and visualize Spotify's playlist data. The objective is to automate the entire process of data extraction, transformation, and loading (ETL) into Snowflake for advanced analytics and insights using Power BI.

Automation is achieved using AWS CloudWatch triggers, Lambda functions, S3 for storage, and Snowpipe in Snowflake for seamless data ingestion. The pipeline generates actionable insights into song trends, artist performance, and release patterns.

2 Literature Review

Existing research highlights the benefits of automation in data pipelines for reducing manual intervention and ensuring scalability. However, there is limited work integrating serverless architectures with data warehouses like Snowflake. By combining AWS services, Snowflake, and Power BI, this project demonstrates an efficient and scalable solution for handling large datasets.

3 Methodology

The methodology involves three phases: automated data collection, transformation, and visualization, ensuring a fully serverless and scalable implementation.

3.1 Data Collection

Spotify API data was extracted every 8 hours using AWS Lambda functions triggered by CloudWatch. Raw JSON data was stored in an S3 bucket, ensuring real-time availability for downstream processing.

3.2 Data Preprocessing

S3 event triggers invoked a second Lambda function to normalize the data into three entities: 'Songs', 'Artists', and 'Albums'. The normalized data was stored as CSV files in S3. Snowpipe automatically ingested the transformed data into Snowflake tables, further automating the pipeline.

3.3 Analysis Techniques

SQL views in Snowflake were used to aggregate and organize data for analysis. Power BI dashboards provided interactive visualizations. Additionally, a machine learning model trained on Snowflake data was used to predict song popularity based on features like duration and artist attributes.

4 Results

Key results from the project include:

- Automation: All processes from data extraction to ingestion into Snowflake were automated using CloudWatch, event triggers, and Snowpipe.
- ETL Pipeline: Successfully extracted, transformed, and loaded data into Snowflake for analytical processing.
- Power BI Dashboards: Interactive insights into artist performance, song trends, and release timing.
- Machine Learning Model: Predicted song popularity, aiding playlist optimization.

5 Discussion

The project demonstrates the benefits of serverless architecture for building scalable, automated pipelines. Challenges included maintaining data quality and integrating machine learning into the ETL process. These were addressed through event-driven architectures and Snowflake's querying capabilities.

6 Conclusion

This project successfully implemented a serverless ETL pipeline integrating AWS, Snowflake, and Power BI for Spotify playlist analytics. The automated workflow reduced manual intervention, and the addition of a machine learning model enhanced playlist optimization. Future enhancements include broader playlist analysis and embedding predictive analytics directly into dashboards.

7 References

References

A Appendix A: Code

```
1 # Spotify Data Transformation Lambda Function
2 def lambda_handler(event, context):
3     playlist_links = ["https://open.spotify.com/playlist/37i9dQZF1DXcBWIGoYBM5M"]
4     raw_data = extract_spotify_data(playlist_links)
5     normalized_data = normalize_data(raw_data)
6     upload_to_s3(normalized_data, 'spotify-etl-pipeline-sn/transform_data/')
```

Listing 1: Lambda Function for Transformation

```
1 # Snowflake/SpotifyPlaylistPowerBIViews.sql
2
3 CREATE OR REPLACE VIEW SpotifyPlaylist_DW.Views.top_5_artists_view AS
4 WITH DistinctSongs AS (
5     SELECT DISTINCT song_id, popularity, artist_id
6     FROM SpotifyPlaylist_DW.SpotifyPlaylist_DW_SCHEMA.song
```

```

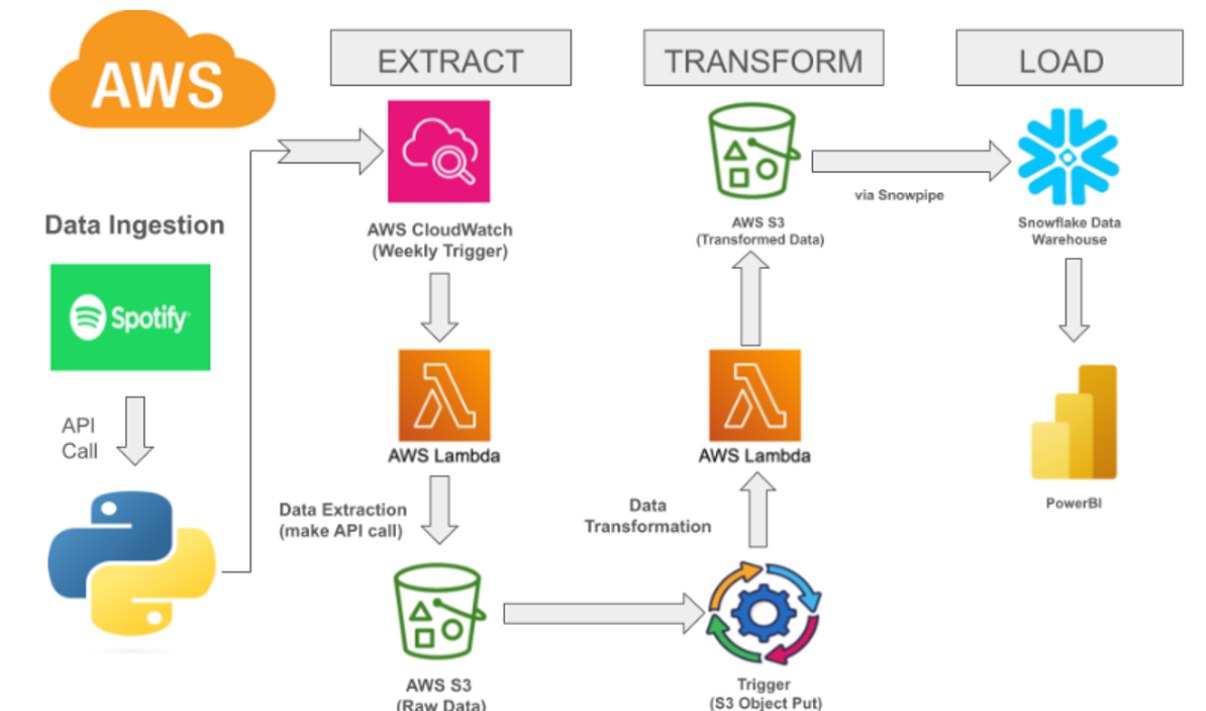
7 )
8 SELECT
9     artist.artist_name,
10     COUNT(song.song_id) AS song_count
11 FROM
12     DistinctSongs AS song
13 JOIN
14     SpotifyPlaylist_DW.SpotifyPlaylist_DW_SCHEMA.artist AS artist
15     ON song.artist_id = artist.artist_id
16 WHERE
17     song.song_id IN (
18         SELECT song_id
19         FROM DistinctSongs
20         ORDER BY popularity DESC
21         LIMIT 50
22     )
23 GROUP BY
24     artist.artist_name
25 ORDER BY
26     song_count DESC
27 LIMIT 5;
28
29 SELECT * FROM SpotifyPlaylist_DW.Views.top_5_artists_view;

```

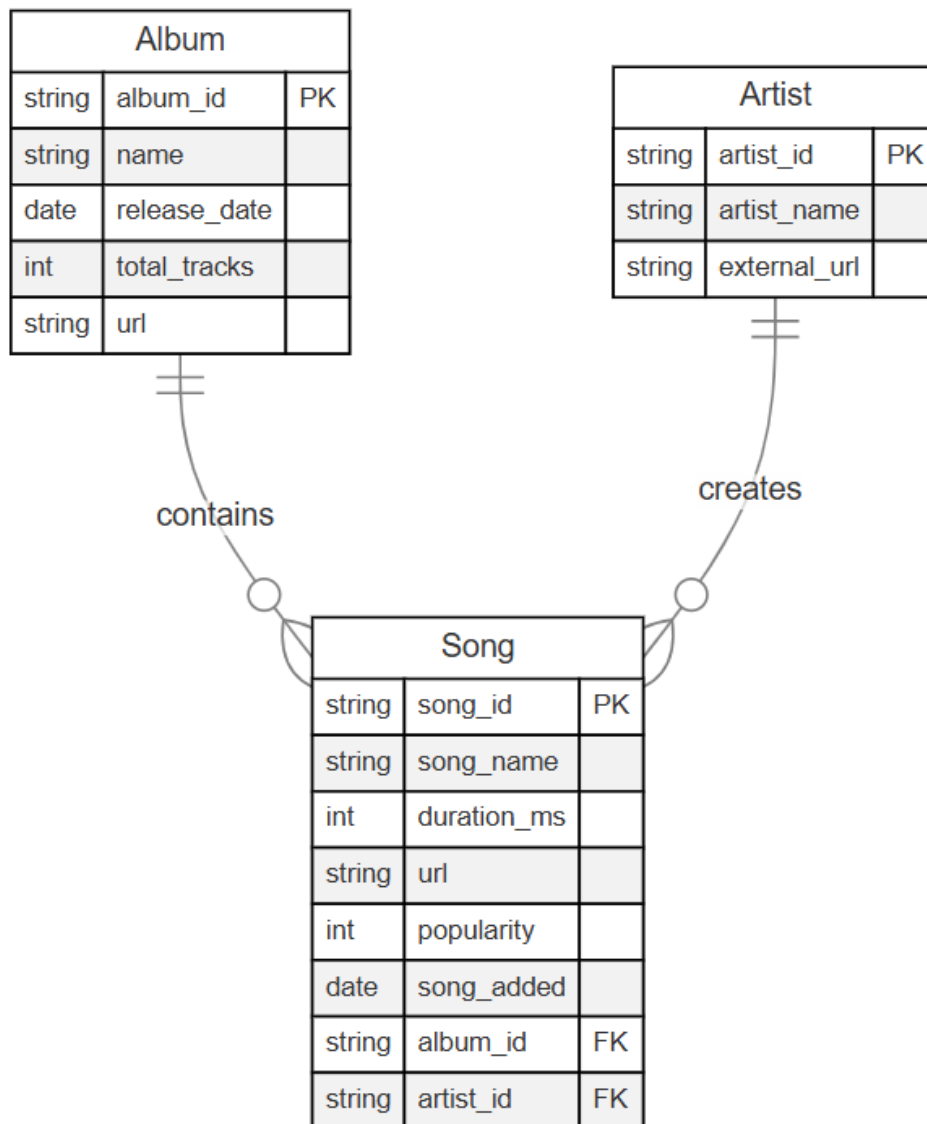
Listing 2: SpotifyPlaylist PowerBI SQL Views

B Appendix B: Additional Figures

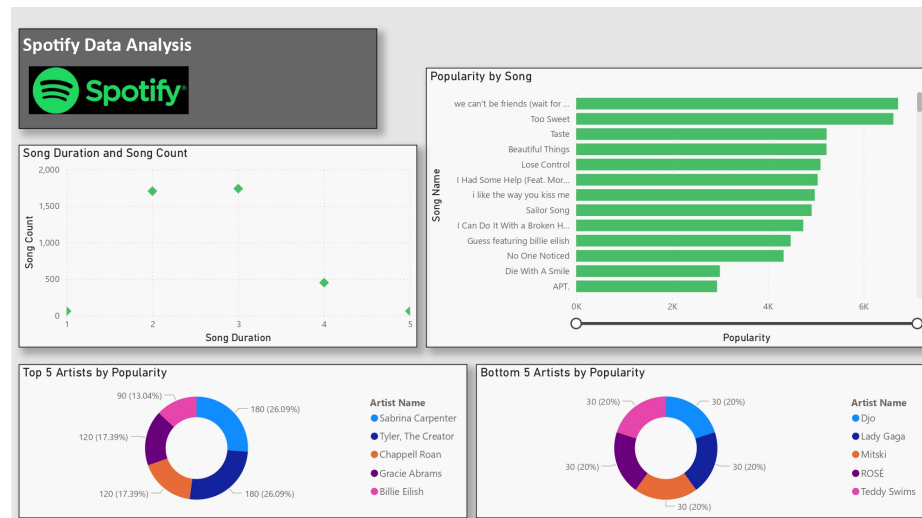
B.1 Pipeline Architecture



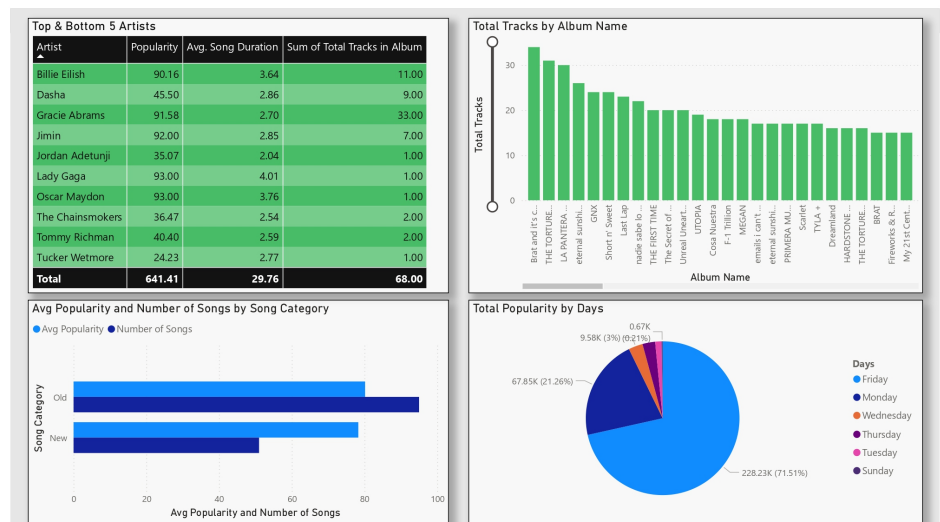
B.2 Data Model



B.3 Power BI Dashboards



- Artist performance metrics:

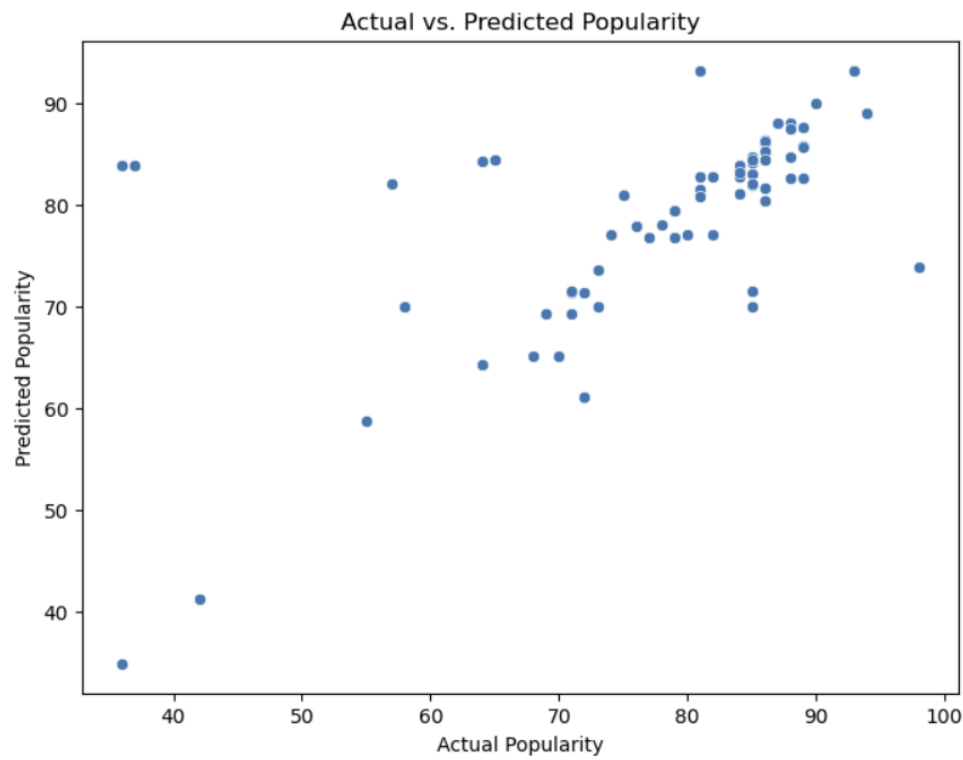


- Song popularity trends:

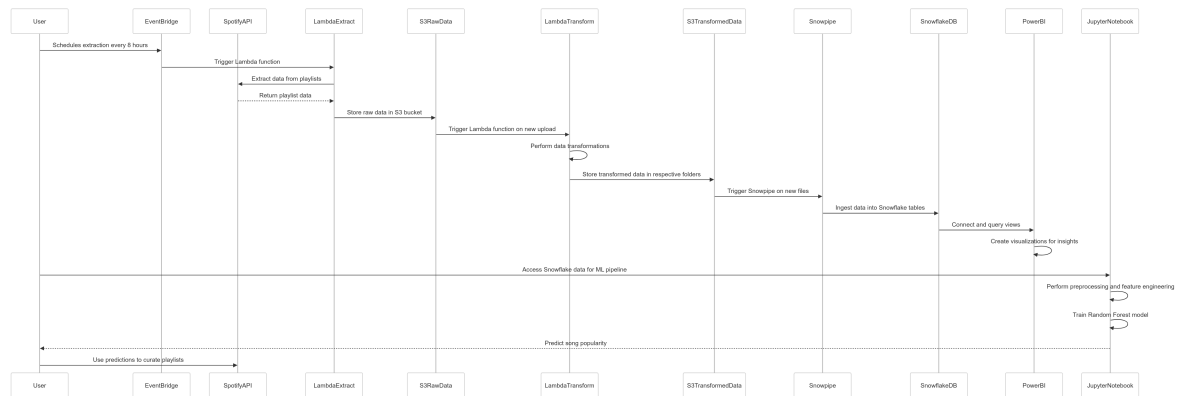


- Release timing analysis:

B.4 Machine Learning Model



B.5 Activity Diagram



B.6 Class Diagram

